

# FaceNet

A Unified Embedding for Face Recognition and Clustering

**The mean squares**

Param Pujara 2020202008

Sanyam Jain 2020201006

Shubham Singh 2020202002

# Problem Statement

Despite significant recent advances in the field of face recognition, implementing face verification and recognition efficiently at scale presents serious challenges to current approaches.

# Solution Approach

- A system is described which converts an image to a fixed size embedding.
- Distances between any two embeddings is a measure of the similarity of the faces.
- The embedding is produced by training a siamese network with triplet loss.

# Introduction to FaceNet

- A unified system for face verification (is this the same person), recognition (who is this person) and clustering (find common people among these faces).
- The embedding generated is an Euclidean embedding per image using a deep convolutional network.
- The network is trained such that the squared L2 distances in the embedding space directly correspond to face similarity.
- Faces of the same person have small distances and faces of distinct people have large distances.
- Face verification simply involves thresholding the distance between the two embeddings, recognition becomes a k-NN classification problem; and clustering can be achieved using off-the-shelf techniques such as k-means or agglomerative clustering.
- Dataset to be used for training:- Labeled Faces in the Wild (LFW).

# Dataset Description

- Database of face photographs designed for studying the problem of unconstrained face recognition.
- For training, we are considering a subset of classes with 3 images each. That comes up to be 610 labels.
- All the images from LFW dataset are split into 2 sets -
  - For validation, to find appropriate threshold for distances of similar faces and different faces.
  - For testing, the images we are taking all labels in the LFW dataset and performing verification/recognition tasks.

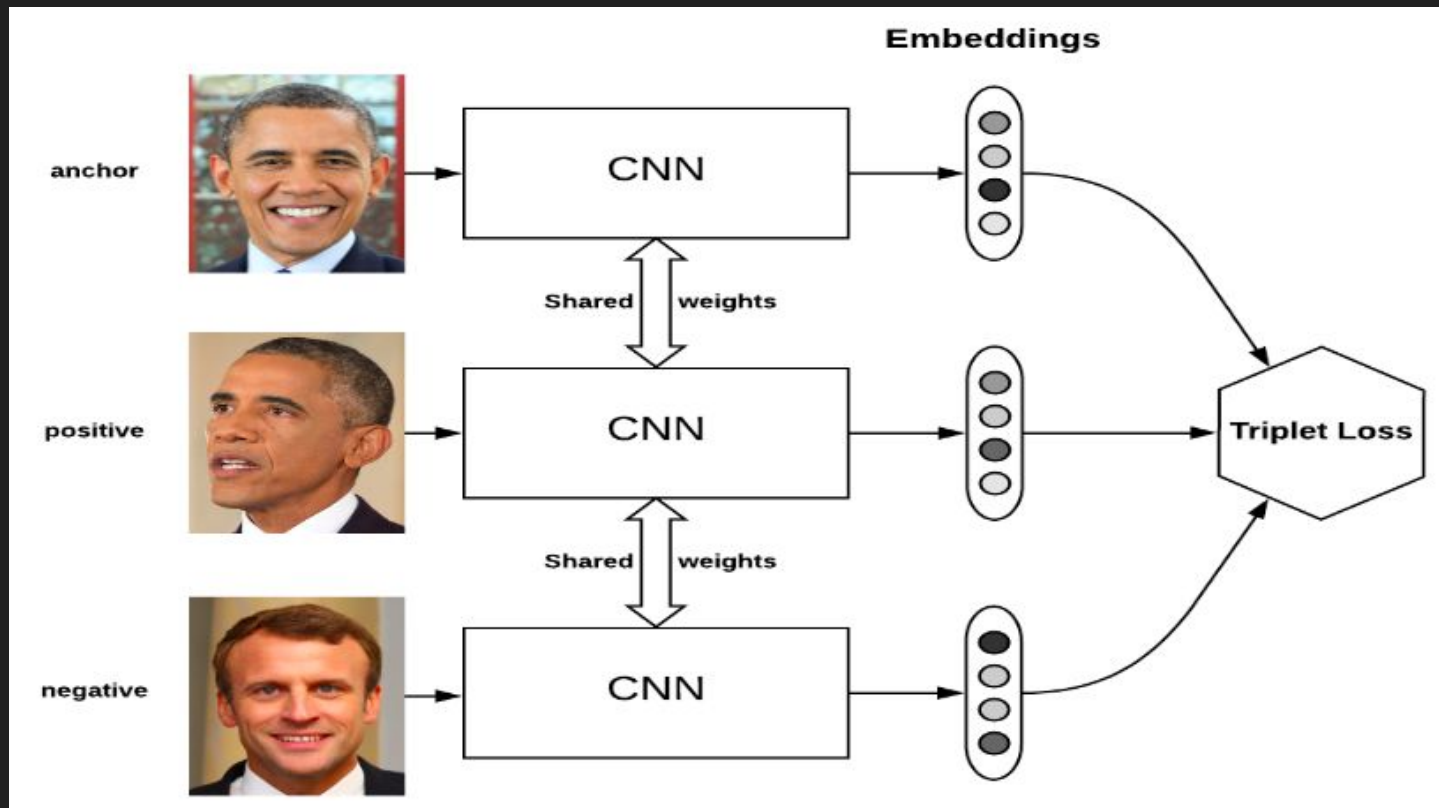
# Expected Results

- The model will produce a fixed size embedding for a test data.
- If the test data belongs to a particular label then the distance between the embedding generated and stored embedding for that class should be less than a threshold.

# Siamese Networks

- Siamese Networks are basically two same neural networks, the output of which are passed through the a function that calculates the distance between the inputs.
- For inputs belonging to same class, the outputs are close (i.e., distance is less) and for inputs belonging to different classes, the distance between outputs should be large.
- Triplet loss can be used to train a siamese network.

# Siamese Network(Cont..)






# Triplet loss

- It is a distance based loss function that operates on three inputs:
  1. anchor (a) : is any arbitrary data point,
  2. positive (p) : which is the same class as the anchor
  3. and negative (n) : which is a different class from the anchor
- Mathematically, it is defined as:  $L = \max(d(a,p) - d(a,n) + \text{margin}, 0)$ .
- We minimize this loss, which pushes  $d(a,p)$  to 0 and  $d(a,n)$  to be greater than  $d(a,p) + \text{margin}$ .
- After the training, the positive examples will be closer to the anchor while the negative examples will be farther from it.

# Triplet Loss(Cont..)



# Triplet loss - loss function

$$\sum_i^N \left[ \|f(x_i^a) - f(x_i^p)\|_2^2 - \|f(x_i^a) - f(x_i^n)\|_2^2 + \alpha \right]_+$$

# Model Summary

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 64)	1792
conv2d_1 (Conv2D)	(None, 15, 15, 128)	73856
conv2d_2 (Conv2D)	(None, 8, 8, 64)	73792
conv2d_3 (Conv2D)	(None, 4, 4, 64)	4160
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
Total params: 842,624		
Trainable params: 842,624		
Non-trainable params: 0		

Model: "functional\_1"


Layer (type)	Output Shape	Param #	Connected to
anchor_input (InputLayer)	[(None, 60, 60, 3)]	0	
positive_input (InputLayer)	[(None, 60, 60, 3)]	0	
negative_input (InputLayer)	[(None, 60, 60, 3)]	0	
sequential (Sequential)	(None, 128)	842624	anchor_input[0][0] positive_input[0][0] negative_input[0][0]
merged_layer (Concatenate)	(None, 384)	0	sequential[0][0] sequential[1][0] sequential[2][0]

Total params: 842,624  
Trainable params: 842,624  
Non-trainable params: 0


# Training and Validation

- For training and validation, we have considered a small subset of the original dataset.
- For validation accuracy, we have used K-NN for better accuracy.

# Epoch vs Loss



# Scatter Plot for validation data and clusters



# Validation Accuracy

K	Accuracy
1	83.33%
3	83.33%
5	79.16%
7	70.83%



# Code Snippets

## Triplet Generation

```
1 def generate_triplets(x, y, num_same = 4, num_diff = 4):
2     anchor_images = np.array([]).reshape((-1,) + x.shape[1:])
3     same_images = np.array([]).reshape((-1,) + x.shape[1:])
4     diff_images = np.array([]).reshape((-1,) + x.shape[1:])
5
6     for i in range(len(y)):
7         point = y[i]
8         anchor = x[i]
9
10        same_pairs = np.where(y == point)[0]
11        same_pairs = np.delete(same_pairs, np.where(same_pairs == i))
12        diff_pairs = np.where(y != point)[0]
13
14        same = x[np.random.choice(same_pairs, num_same)]
15        diff = x[np.random.choice(diff_pairs, num_diff)]
16
17        anchor_images = np.concatenate((anchor_images, np.tile(anchor, (num_same * num_diff, 1, 1, 1))), axis = 0)
18
19        for s in same:
20            same_images = np.concatenate((same_images, np.tile(s, (num_same, 1, 1, 1))), axis = 0)
21
22        diff_images = np.concatenate((diff_images, np.tile(diff, (num_diff, 1, 1, 1))), axis = 0)
23
24    return anchor_images, same_images, diff_images
```

# Code Snippets

## Triplet Loss Function

```
1 def triplet_loss(y_true, y_pred, alpha = 0.2):
2     total_length = y_pred.shape.as_list()[-1]
3     anchor, positive, negative = y_pred[:, :int(1/3*total_length)], \
4         y_pred[:, int(1/3*total_length):int(2/3*total_length)], y_pred[:, int(2/3*total_length):]
5
6     pos_dist = tf.reduce_sum(tf.square(anchor - positive), axis=-1)
7     neg_dist = tf.reduce_sum(tf.square(anchor - negative), axis=-1)
8     basic_loss = pos_dist - neg_dist + alpha
9     loss = tf.reduce_sum(tf.maximum(basic_loss, 0.0))
10    return loss
```

# Models and experiments

- We have tried different deep learning structures and experimented with them.
  - Experiments on Batch Normalisation
  - Max/Min/Avg pooling
  - Dropout
  - optimizers
  - epochs and different batch sizes
- We tried different values of K in KNN classification.
- We trained our model to different numbers of triplets.
- We tried different subsets of LFW dataset for training model.

# Accuracy on LFW

Model	Accuracy
Margin: 0.4, K=7, w/o dropout	52.70%
Margin: 0.4, K=9, with dropout	51.45%
Margin: 0.5, k = 1, w/o dropout	55.96%
Margin: 1, k = 7, w/o dropout, Batch normalization	52.74%

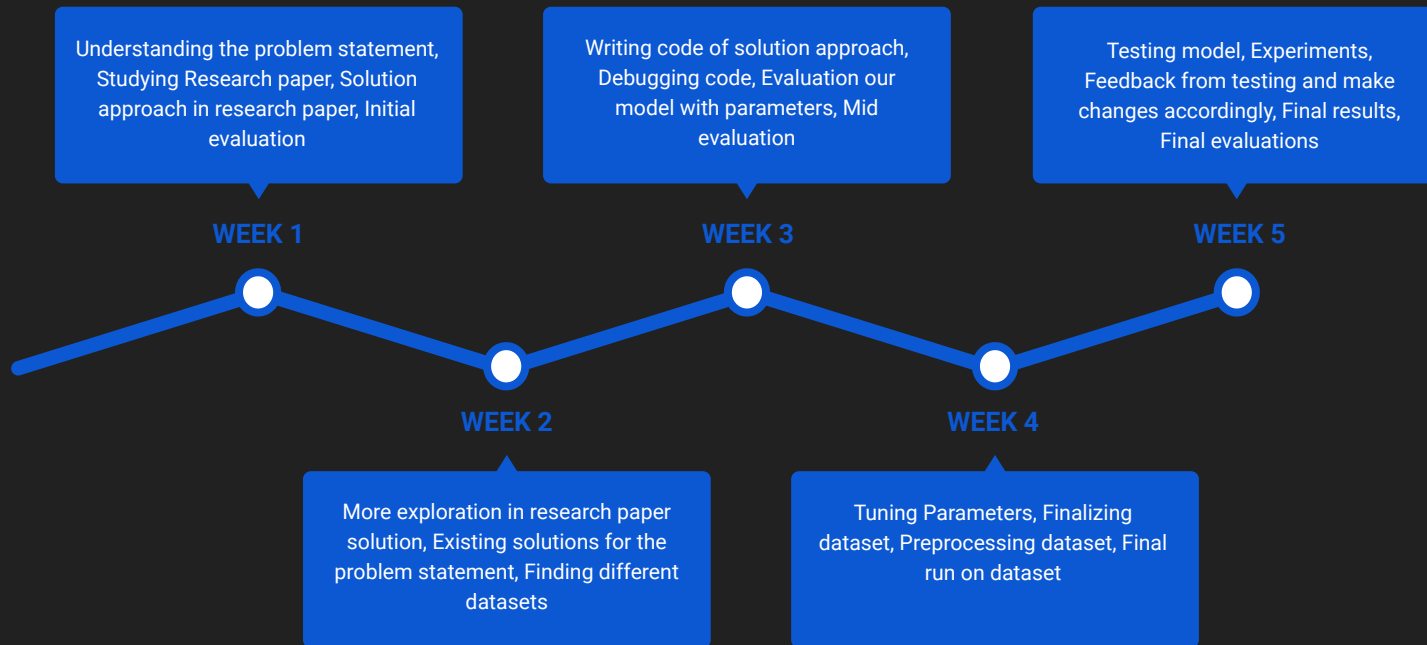
# Final Model & Accuracy

Margin: 0.6, k=9, w/o dropout

**Accuracy = 61.32%**

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_1 (Conv2D)	(None, 15, 15, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 128)	0
conv2d_2 (Conv2D)	(None, 8, 8, 64)	73792
conv2d_3 (Conv2D)	(None, 4, 4, 64)	4160
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 128)	32896
Total params: 842,624		
Trainable params: 842,624		
Non-trainable params: 0		

# Timeline



# Individual contribution

- Data Preprocessing & face detection - Param and Sanyam
- Generating triplets - Shubham and Sanyam
- Siamese Network and Triplet Loss function - Param and Shubham
- Experimenting with Different parameters - All

# References

1. <https://arxiv.org/abs/1503.03832>
2. <https://machinelearningmastery.com/how-to-develop-a-face-recognition-system-using-facenet-in-keras-and-an-svm-classifier/>
3. <https://towardsdatascience.com/siamese-network-triplet-loss-b4ca82c1aec8>
4. <https://towardsdatascience.com/understand-and-implement-resnet-50-with-tensorflow-2-0-1190b9b52691>
5. <https://www.coursera.org/lecture/convolutional-neural-networks/triplet-loss-HuUtN>
6. <https://medium.com/analytics-vidhya/facenet-architecture-part-1-a062d5d918a1>
7. <https://medium.com/@anilmatcha/summary-for-facenet-a-unified-embedding-for-face-recognition-and-clustering-cfe878027a3d>